


```

In [5]: %%writefile test_helper.py
# A testing helper
#https://pypi.python.org/pypi/test_helper/0.2
import hashlib

class TestFailure(Exception):
    pass
class PrivateTestFailure(Exception):
    pass

class Test(object):
    passed = 0
    numTests = 0
    failFast = False
    private = False

    @classmethod
    def setFailFast(cls):
        cls.failFast = True

    @classmethod
    def setPrivateMode(cls):
        cls.private = True

    @classmethod
    def assertTrue(cls, result, msg=""):
        cls.numTests += 1
        if result == True:
            cls.passed += 1
            print "1 test passed."
        else:
            print "1 test failed. " + msg
            if cls.failFast:
                if cls.private:
                    raise PrivateTestFailure(msg)
                else:
                    raise TestFailure(msg)

    @classmethod
    def assertEquals(cls, var, val, msg=""):
        cls.assertTrue(var == val, msg)

    @classmethod
    def assertEqualsHashed(cls, var, hashed_val, msg=""):
        cls.assertEquals(cls._hash(var), hashed_val, msg)

    @classmethod
    def printStats(cls):
        print "{0} / {1} test(s) passed.".format(cls.passed, cls.numTests)

    @classmethod
    def _hash(cls, x):
        return hashlib.sha1(str(x)).hexdigest()

```

Writing test_helper.py

```
In [6]: # TEST One-hot-encoding (1a)
from test_helper import Test

Test.assertEqualHashed(sampleOHEDictManual[(0, 'bear')],
                        'b6589fc6ab0dc82cf12099d1c2d40ab994e8410c',
                        "incorrect value for sampleOHEDictManual[(0, 'bear')]",
Test.assertEqualHashed(sampleOHEDictManual[(0, 'cat')],
                        '356a192b7913b04c54574d18c28d46e6395428ab',
                        "incorrect value for sampleOHEDictManual[(0, 'cat')]",
Test.assertEqualHashed(sampleOHEDictManual[(0, 'mouse')],
                        'da4b9237baccdf19c0760cab7aec4a8359010b0',
                        "incorrect value for sampleOHEDictManual[(0, 'mouse')]",
Test.assertEqualHashed(sampleOHEDictManual[(1, 'black')],
                        '77de68daecd823babbb58edb1c8e14d7106e83bb',
                        "incorrect value for sampleOHEDictManual[(1, 'black')]",
Test.assertEqualHashed(sampleOHEDictManual[(1, 'tabby')],
                        '1b6453892473a467d07372d45eb05abc2031647a',
                        "incorrect value for sampleOHEDictManual[(1, 'tabby')]",
Test.assertEqualHashed(sampleOHEDictManual[(2, 'mouse')],
                        'ac3478d69a3c81fa62e60f5c3696165a4e5e6ac4',
                        "incorrect value for sampleOHEDictManual[(2, 'mouse')]",
Test.assertEqualHashed(sampleOHEDictManual[(2, 'salmon')],
                        'c1dfd96eea8cc2b62785275bca38ac261256e278',
                        "incorrect value for sampleOHEDictManual[(2, 'salmon')]",
Test.assertEqual(len(sampleOHEDictManual.keys()), 7,
                  "incorrect number of keys in sampleOHEDictManual")

1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

1b

```
In [7]: import numpy as np
from test_helper import Test
```

```
In [8]: # TODO: Replace <FILL IN> with appropriate code
aDense = np.array([0., 3., 0., 4.])
aSparse = SparseVector(4, [1,3], [3,4]) #<FILL IN>

bDense = np.array([0., 0., 0., 1.])
bSparse = SparseVector(4, [3], [1]) #<FILL IN>

w = np.array([0.4, 3.1, -1.4, -.5])
print aDense.dot(w)
print aSparse.dot(w)
print bDense.dot(w)
print bSparse.dot(w)
```

7.3
7.3
-0.5
-0.5

```
In [9]: # TEST Sparse Vectors (1b)
Test.assertTrue(isinstance(aSparse, SparseVector), 'aSparse needs to be an
Test.assertTrue(isinstance(bSparse, SparseVector), 'aSparse needs to be an
Test.assertTrue(aDense.dot(w) == aSparse.dot(w),
                'dot product of aDense and w should equal dot product of a
Test.assertTrue(bDense.dot(w) == bSparse.dot(w),
                'dot product of bDense and w should equal dot product of b
```

1 test passed.
1 test passed.
1 test passed.
1 test passed.

1c

```
In [10]: # TODO: Replace <FILL IN> with appropriate code
sampleOneOHEFeatManual = SparseVector(7, [2,3], [1,1]) #<FILL IN>
sampleTwoOHEFeatManual = SparseVector(7, [1,4,5], [1,1,1]) #<FILL IN>
sampleThreeOHEFeatManual = SparseVector(7, [0,2,6], [1,1,1]) #<FILL IN>
```

```
In [11]: # TEST OHE Features as sparse vectors (1c)
Test.assertTrue(isinstance(sampleOneOHEFeatManual, SparseVector),
                  'sampleOneOHEFeatManual needs to be a SparseVector')
Test.assertTrue(isinstance(sampleTwoOHEFeatManual, SparseVector),
                  'sampleTwoOHEFeatManual needs to be a SparseVector')
Test.assertTrue(isinstance(sampleThreeOHEFeatManual, SparseVector),
                  'sampleThreeOHEFeatManual needs to be a SparseVector')
Test.assertEqualHashed(sampleOneOHEFeatManual,
                        'ecc00223d141b7bd0913d52377cee2cf5783abd6',
                        'incorrect value for sampleOneOHEFeatManual')
Test.assertEqualHashed(sampleTwoOHEFeatManual,
                        '26b023f4109e3b8ab32241938e2e9b9e9d62720a',
                        'incorrect value for sampleTwoOHEFeatManual')
Test.assertEqualHashed(sampleThreeOHEFeatManual,
                        'c04134fd603ae115395b29dcabe9d0c66fbdc8a7',
                        'incorrect value for sampleThreeOHEFeatManual')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

1d

```
In [12]: # TODO: Replace <FILL IN> with appropriate code
def oneHotEncoding(rawFeats, OHEDict, numOHEFeats):
    """Produce a one-hot-encoding from a list of features and an OHE dicti

    Note:
        You should ensure that the indices used to create a SparseVector a

    Args:
        rawFeats (list of (int, str)): The features corresponding to a sin
            feature consists of a tuple of featureID and the feature's val
        OHEDict (dict): A mapping of (featureID, value) to unique integer.
        numOHEFeats (int): The total number of unique OHE features (combin
            value).

    Returns:
        SparseVector: A SparseVector of length numOHEFeats with indicies e
            identifiers for the (featureID, value) combinations that occur
            with values equal to 1.0.
    """
    #<FILL IN>
    # build index array
    ids = [OHEDict[x] for x in rawFeats]
    ones = [1]*len(ids)
    return SparseVector(numOHEFeats, np.sort(ids), ones)

# Calculate the number of features in sampleOHEDictManual
numSampleOHEFeats = len(sampleOHEDictManual) #<FILL IN>

# Run oneHotEnoding on sampleOne
sampleOneOHEFeat = oneHotEncoding(sampleOne, sampleOHEDictManual, numSampl

print sampleOneOHEFeat

(7, [2,3], [1.0,1.0])
```

```
In [13]: # TEST Define an OHE Function (1d)
Test.assertTrue(sampleOneOHEFeat == sampleOneOHEFeatManual,
                'sampleOneOHEFeat should equal sampleOneOHEFeatManual')
Test.assertEquals(sampleOneOHEFeat, SparseVector(7, [2,3], [1.0,1.0]),
                'incorrect value for sampleOneOHEFeat')
Test.assertEquals(oneHotEncoding([(1, 'black'), (0, 'mouse')], sampleOHEDi
                numSampleOHEFeats), SparseVector(7, [2,3]
                'incorrect definition for oneHotEncoding')

1 test passed.
1 test passed.
1 test passed.
```

1e

```
In [14]: # TODO: Replace <FILL IN> with appropriate code
sampleOHEData = sampleDataRDD.map(lambda p: oneHotEncoding(p, sampleOHEDic
print sampleOHEData.collect())

[ SparseVector(7, {2: 1.0, 3: 1.0}), SparseVector(7, {1: 1.0, 4: 1.0, 5:
  1.0}), SparseVector(7, {0: 1.0, 3: 1.0, 6: 1.0}) ]
```

```
In [15]: # TEST Apply OHE to a dataset (1e)
sampleOHEDataValues = sampleOHEData.collect()
Test.assertTrue(len(sampleOHEDataValues) == 3, 'sampleOHEData should have
Test.assertEquals(sampleOHEDataValues[0], SparseVector(7, {2: 1.0, 3: 1.0}
    'incorrect OHE for first sample')
Test.assertEquals(sampleOHEDataValues[1], SparseVector(7, {1: 1.0, 4: 1.0,
    'incorrect OHE for second sample')
Test.assertEquals(sampleOHEDataValues[2], SparseVector(7, {0: 1.0, 3: 1.0,
    'incorrect OHE for third sample')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

Part 2

2a

```
In [16]: # TODO: Replace <FILL IN> with appropriate code
sampleDistinctFeats = sampleDataRDD.flatMap(lambda p: p).distinct().sortBy
# pretty neat it can sort key then value, basically for composite key sort
#<FILL IN>)
```

```
In [17]: # TEST Pair RDD of (featureID, category) (2a)
Test.assertEquals(sorted(sampleDistinctFeats.collect()),
    [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
     (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
    'incorrect value for sampleDistinctFeats')

1 test passed.
```

2b

```
In [18]: # TODO: Replace <FILL IN> with appropriate code
sampleOHEDict = sampleDistinctFeats.zipWithIndex().collectAsMap()
print sampleOHEDict

{(2, 'mouse'): 5, (0, 'cat'): 1, (0, 'bear'): 0, (2, 'salmon'): 6, (1,
'tabby'): 4, (1, 'black'): 3, (0, 'mouse'): 2}
```

```
In [19]: # TEST OHE Dictionary from distinct features (2b)
Test.assertEquals(sorted(sampleOHEDict.keys()),
                  [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
                   (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
                  'sampleOHEDict has unexpected keys')
Test.assertEquals(sorted(sampleOHEDict.values()), range(7), 'sampleOHEDict

1 test passed.
1 test passed.
```

2c

```
In [20]: # TODO: Replace <FILL IN> with appropriate code
def createOneHotDict(inputData):
    """Creates a one-hot-encoder dictionary based on the input data.

    Args:
        inputData (RDD of lists of (int, str)): An RDD of observations whe
            made up of a list of (featureID, value) tuples.

    Returns:
        dict: A dictionary where the keys are (featureID, value) tuples an
            unique integers.
    """
    #<FILL IN>
    return inputData.flatMap(lambda p: p).distinct().sortBy(lambda p: p).z

sampleOHEDictAuto = createOneHotDict(sampleDataRDD) #<FILL IN>
print sampleOHEDictAuto

{(2, 'mouse'): 5, (0, 'cat'): 1, (0, 'bear'): 0, (2, 'salmon'): 6, (1,
'tabby'): 4, (1, 'black'): 3, (0, 'mouse'): 2}
```

```
In [21]: # TEST Automated creation of an OHE dictionary (2c)
Test.assertEquals(sorted(sampleOHEDictAuto.keys()),
                  [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
                   (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
                  'sampleOHEDictAuto has unexpected keys')
Test.assertEquals(sorted(sampleOHEDictAuto.values()), range(7),
                  'sampleOHEDictAuto has unexpected values')

1 test passed.
1 test passed.
```

Part 3


```
In [22]: # Run this code to view Criteo's agreement
        from IPython.lib.display import IFrame

        IFrame("http://labs.criteo.com/downloads/2014-kaggle-display-advertising-c",
                600, 250)
```

Out[22]:

```

In [25]: # TODO: Replace <FILL IN> with appropriate code
# Just replace <FILL IN> with the url for dac_sample.tar.gz
import glob
import os.path
import tarfile
import urllib
import urlparse

# Paste url, url should end with: dac_sample.tar.gz
url = 'http://labs.criteo.com/wp-content/uploads/2015/04/dac_sample.tar.gz'

url = url.strip()
baseDir = os.path.join('home')
inputPath = os.path.join('cloudera', 'dac_sample.txt')
fileName = os.path.join(baseDir, inputPath)
inputDir = os.path.split(fileName)[0]

def extractTar(check = False):
    # Find the zipped archive and extract the dataset
    tars = glob.glob('dac_sample*.tar.gz')
    if check and len(tars) == 0:
        return False

    if len(tars) > 0:
        try:
            tarFile = tarfile.open(tars[0])
        except tarfile.ReadError:
            if not check:
                print 'Unable to open tar.gz file. Check your URL.'
            return False

        tarFile.extract('dac_sample.txt', path=inputDir)
        print 'Successfully extracted: dac_sample.txt'
        return True
    else:
        print 'You need to retry the download with the correct url.'
        print ('Alternatively, you can upload the dac_sample.tar.gz file to the '
              'directory')
        return False

if os.path.isfile(fileName):
    print 'File is already available. Nothing to do.'
elif extractTar(check = True):
    print 'tar.gz file was already available.'
elif not url.endswith('dac_sample.tar.gz'):
    print 'Check your download url. Are you downloading the Sample dataset?'
else:
    # Download the file and store it in the same directory as this notebook
    try:
        urllib.urlretrieve(url, os.path.basename(urlparse.urlsplit(url).path))
    except IOError:
        print 'Unable to download and store: {0}'.format(url)

    extractTar()

```

```
In [33]: import os.path
#baseDir = os.path.join('home')
#inputPath = os.path.join('cloudera', 'dac_sample.txt')#
#fileName = os.path.join(baseDir, inputPath)
fileName = '/home/cloudera/dac_sample.txt'

if os.path.isfile(fileName):
    rawData = (sc
                .textFile(fileName, 2)
                .map(lambda x: x.replace('\t', ','))) # work with either '
    print rawData.take(1)

[u'0,1,1,5,0,1382,4,15,2,181,1,2,,2,68fd1e64,80e26c9b,fb936136,7b4723c4
,25c83c98,7e0ccccf,de7995b8,1f89b562,a73ee510,a8cd5504,b2cb9c98,37c9c16
4,2824a5f6,1adce6ef,8ba8b39a,891b62e7,e5ba7672,f54016b9,21ddcdc9,b1252a
9d,07b5194c,,3a171ecb,c5c50484,e8b83407,9727dd16']
```

3a

```
In [34]: # TODO: Replace <FILL IN> with appropriate code
weights = [.8, .1, .1]
seed = 42
# Use randomSplit with weights and seed
rawTrainData, rawValidationData, rawTestData = rawData.randomSplit(weights)
# Cache the data
#<FILL IN>
rawTrainData.cache()
rawValidationData.cache()
rawTestData.cache()

# count the data
nTrain = rawTrainData.count()
nVal = rawValidationData.count()
nTest = rawTestData.count() #<FILL IN>
print nTrain, nVal, nTest, nTrain + nVal + nTest

79911 10075 10014 100000
[u'0,1,1,5,0,1382,4,15,2,181,1,2,,2,68fd1e64,80e26c9b,fb936136,7b4723c4
,25c83c98,7e0ccccf,de7995b8,1f89b562,a73ee510,a8cd5504,b2cb9c98,37c9c16
4,2824a5f6,1adce6ef,8ba8b39a,891b62e7,e5ba7672,f54016b9,21ddcdc9,b1252a
9d,07b5194c,,3a171ecb,c5c50484,e8b83407,9727dd16']
```

```
In [35]: # TEST Loading and splitting the data (3a)
Test.assertTrue(all([rawTrainData.is_cached, rawValidationData.is_cached,
                    'you must cache the split data']))
Test.assertEquals(nTrain, 79911, 'incorrect value for nTrain')
Test.assertEquals(nVal, 10075, 'incorrect value for nVal')
Test.assertEquals(nTest, 10014, 'incorrect value for nTest')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

3b

```
In [36]: # TODO: Replace <FILL IN> with appropriate code
def parsePoint(point):
    """Converts a comma separated string into a list of (featureID, value)

    Note:
        featureIDs should start at 0 and increase to the number of feature

    Args:
        point (str): A comma separated string where the first value is the
            are features.

    Returns:
        list: A list of (featureID, value) tuples.
    """
    #<FILL IN>
    fea = point.strip().split(',')
    return [(i, fea[i]) for i in range(len(fea))]

parsedTrainFeat = rawTrainData.map(parsePoint)

numCategories = (parsedTrainFeat
                 .flatMap(lambda x: x)
                 .distinct()
                 .map(lambda x: (x[0], 1))
                 .reduceByKey(lambda x, y: x + y)
                 .sortByKey()
                 .collect())

print numCategories[2][1]

# print numCategories[2][1]
855
```

```
In [37]: # TEST Extract features (3b)
Test.assertEquals(numCategories[2][1], 855, 'incorrect implementation of p
Test.assertEquals(numCategories[2][1], 4, 'incorrect implementation of p
1 test passed.
1 test passed.
```

3c

```
In [38]: # TODO: Replace <FILL IN> with appropriate code
ctrOHEDict = createOneHotDict(parsedTrainFeat) #<FILL IN>
numCtrOHEFeats = len(ctrOHEDict.keys())
print numCtrOHEFeats

print ctrOHEDict[(0, 1)]
233286
0
```

```
In [39]: # TEST Create an OHE dictionary from the dataset (3c)
Test.assertEquals(numCtrOHEFeats, 233286, 'incorrect number of features in
Test.assertEquals(0, 1) is not OHEDict. Incorrect features in OHEDict')
1 test passed.
1 test passed.
```

3d

```
In [40]: from pyspark.mllib.regression import LabeledPoint

# TODO: Replace <FILL IN> with appropriate code
def parseOHEPoint(point, OHEDict, numOHEFeats):
    """Obtain the label and feature vector for this raw observation.

    Note:
        You must use the function `oneHotEncoding` in this implementation
        of this lab may not function as expected.

    Args:
        point (str): A comma separated string where the first value is the
            are features.
        OHEDict (dict of (int, str) to int): Mapping of (featureID, value)
        numOHEFeats (int): The number of unique features in the training d

    Returns:
        LabeledPoint: Contains the label for the observation and the one-h
            raw features based on the provided OHE dictionary.
    """
    #<FILL IN>
    label, fea = point.strip().split(',', 1)
    return LabeledPoint(label, oneHotEncoding(parsePoint(point), OHEDict,

OHETrainData = rawTrainData.map(lambda point: parseOHEPoint(point, ctrOHED
OHETrainData.cache()
print OHETrainData.take(1)

# Check that oneHotEncoding function was used in parseOHEPoint
backupOneHot = oneHotEncoding
oneHotEncoding = None
withOneHot = False
try: parseOHEPoint(rawTrainData.take(1)[0], ctrOHEDict, numCtrOHEFeats)
except TypeError: withOneHot = True
oneHotEncoding = backupOneHot

[LabeledPoint(0.0, (233286, [2,147,3206,3467,6199,25073,25873,26398,2707
9,28303,28323,28390,28554,28906,29422,65104,75996,87062,87178,93426,944
29,94641,101154,107147,115594,142624,144958,147589,165315,178594,180856
,181114,182200,183160,214088,214101,222510,224951,229867], [1.0,1.0,1.0,
1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0
,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.0,1.
0]))]
```

```
In [41]: # TEST Apply OHE to the dataset (3d)
numNZ = sum(parsedTrainFeat.map(lambda x: len(x)).take(5))
numNZAlt = sum(OHETrainData.map(lambda lp: len(lp.features.indices)).take(
Test.assertEquals(numNZ, numNZAlt, 'incorrect implementation of parseOHEPo
Test.assertEquals(withOneHot, lenOneHotEncoding, 'not present in parseOHEPoint'))
1 test passed.
1 test passed.
```

```
In [42]: def bucketFeatByCount(featsCount):
    """Bucket the counts by powers of two."""
    for i in range(11):
        size = 2 ** i
        if featsCount <= size:
            return size
    return -1

featsCounts = (OHETrainData
    .flatMap(lambda lp: lp.features.indices)
    .map(lambda x: (x, 1))
    .reduceByKey(lambda x, y: x + y))
featsCountsBuckets = (featsCounts
    .map(lambda x: (bucketFeatByCount(x[1]), 1))
    .filter(lambda (k, v): k != -1)
    .reduceByKey(lambda x, y: x + y)
    .collect())

print(featsCountsBuckets)


[(512, 414), (1024, 255), (2, 24076), (4, 16639), (32, 4755), (64, 2627
), (128, 1476), (256, 748), (16, 7752), (8, 11440), (1, 162813)]
```

```
In [43]: import matplotlib.pyplot as plt
%matplotlib inline

x, y = zip(*featCountsBuckets)
x, y = np.log(x), np.log(y)

def preparePlot(xticks, yticks, figsize=(10.5, 6), hideLabels=False, gridC
               gridWidth=1.0):
    """Template for generating the plot layout."""
    plt.close()
    fig, ax = plt.subplots(figsize=figsize, facecolor='white', edgecolor='
    ax.axes.tick_params(labelcolor='#999999', labelsizes='10')
    for axis, ticks in [(ax.get_xaxis(), xticks), (ax.get_yaxis(), yticks)]:
        axis.set_ticks_position('none')
        axis.set_ticks(ticks)
        axis.label.set_color('#999999')
        if hideLabels: axis.set_ticklabels([])
    plt.grid(color=gridColor, linewidth=gridWidth, linestyle='-')
    map(lambda position: ax.spines[position].set_visible(False), ['bottom'
    return fig, ax

# generate layout and plot data
fig, ax = preparePlot(np.arange(0, 10, 1), np.arange(4, 14, 2))
ax.set_xlabel(r'$\log_e(bucketSize)$'), ax.set_ylabel(r'$\log_e(countInBuc
plt.scatter(x, y, s=14**2, c='#d6ebf2', edgecolors='#8cbfd0', alpha=0.75)
#pass
plt.show()
```



3e

```
In [45]: # TEST Handling unseen features (3e)
numNZVal = (OHEValidationData
             .map(lambda lp: len(lp.features.indices))
             .sum())
Test count Female (numNZVal=1 27000) -> incorrect number of features 1
1 test passed.
```

4a


```
In [46]: from pyspark.mllib.classification import LogisticRegressionWithSGD

# fixed hyperparameters
numIters = 50
stepSize = 10.
regParam = 1e-6
regType = 'l2'
```

```
In [47]: # TODO: Replace <FILL IN> with appropriate code
model0 = LogisticRegressionWithSGD.train(OHETrainData,
                                         iterations=numIters,
                                         step=stepSize,
                                         regParam=regParam,
                                         regType=regType,
                                         intercept=includeIntercept) #<FILL IN>

sortedWeights = sorted(model0.weights)
print sortedWeights[0:5], model0.intercept

/home/cloudera/Downloads/spark-2.0.1-bin-hadoop2.6/python/pyspark/mllib
/classification.py:313: UserWarning: Deprecated in 2.0.0. Use ml.classi
fication.LogisticRegression or LogisticRegressionWithLBFGS.
  "Deprecated in 2.0.0. Use ml.classification.LogisticRegression or "

[-0.45899236853575626, -0.37973707648623972, -0.36996558266753299, -0.3
6934962879928268, -0.32697945415010637] 0.56455084025
```

```
In [48]: # TEST Logistic regression (4a)
Test.assertTrue(np.allclose(model0.intercept, 0.56455084025), 'incorrect')
Test.assertTrue(np.allclose(sortedWeights[0:5],
                             [-0.45899236853575609, -0.37973707648623956, -0.3699655826
                             -0.36934962879928263, -0.32697945415010637])), 'incorrect')

1 test passed.
1 test passed.
```

4b

```
In [49]: # TODO: Replace <FILL IN> with appropriate code
from math import log

def computeLogLoss(p, y):
    """Calculates the value of log loss for a given probability and label.

    Note:
        log(0) is undefined, so when p is 0 we need to add a small value (epsilon)
        and when p is 1 we need to subtract a small value (epsilon) from it.

    Args:
        p (float): A probability between 0 and 1.
        y (int): A label. Takes on the values 0 and 1.

    Returns:
        float: The log loss value.
    """
    epsilon = 10e-12
    #<FILL IN>
    return -log(p+epsilon) if y==1 else -log(1-p+epsilon)

print computeLogLoss(.5, 1)
print computeLogLoss(.5, 0)
print computeLogLoss(.99, 1)
print computeLogLoss(.99, 0)
print computeLogLoss(.01, 1)
print computeLogLoss(.01, 0)
print computeLogLoss(0, 1)
print computeLogLoss(1, 1)
print computeLogLoss(1, 0)

0.69314718054
0.69314718054
0.0100503358434
4.60517018499
4.60517018499
0.0100503358434
25.3284360229
-1.00000008274e-11
25.3284360229
```

```
In [50]: # TEST Log loss (4b)
Test.assertTrue(np.allclose([computeLogLoss(.5, 1), computeLogLoss(.01, 0),
                             computeLogLoss(.99, 1), computeLogLoss(.99, 0),
                             computeLogLoss(.01, 1), computeLogLoss(.01, 0),
                             computeLogLoss(0, 1), computeLogLoss(1, 1),
                             computeLogLoss(1, 0)],
                             [0.69314718056, 0.0100503358535, 4.60517018599,
                              4.60517018599, 0.0100503358434, 25.3284360229,
                              1.00000008275e-11, 25.3284360229,
                              -1.00000008275e-11],
                             atol=1e-11, rtol=1e-11,
                             msg='computeLogLoss is not correct'))
Test.assertTrue(np.allclose([computeLogLoss(0, 1), computeLogLoss(1, 1),
                             computeLogLoss(1, 0)],
                             [25.3284360229, 1.00000008275e-11, 25.3284360229],
                             atol=1e-11, rtol=1e-11,
                             msg='computeLogLoss needs to bound p away from 0 and 1 by epsilon'))

1 test passed.
1 test passed.
```

4c

```
In [51]: # TODO: Replace <FILL IN> with appropriate code
# Note that our dataset has a very high click-through rate by design
# In practice click-through rate can be one to two orders of magnitude low
classOneFracTrain = OHETrainData.map(lambda p: p.label).mean() #<FILL IN>
print classOneFracTrain

logLossTrBase = OHETrainData.map(lambda p: computeLogLoss(classOneFracTrain, p.label))
print 'Baseline Train Logloss = {0:.3f}\n'.format(logLossTrBase)

0.22717773523
Baseline Train Logloss = 0.536
```

```
In [52]: # TEST Baseline log loss (4c)
Test.assertTrue(np.allclose(classOneFracTrain, 0.22717773523), 'incorrect
Test.assertTrue(np.allclose(logLossTrBase, 0.535844), 'incorrect value for

1 test passed.
1 test passed.
```

4d

```
In [53]: # TODO: Replace <FILL IN> with appropriate code
from math import exp #  $\exp(-t) = e^{-t}$ 

def getP(x, w, intercept):
    """Calculate the probability for an observation given a set of weights

    Note:
        We'll bound our raw prediction between 20 and -20 for numerical pu

    Args:
        x (SparseVector): A vector with values of 1.0 for features that ex
            observation and 0.0 otherwise.
        w (DenseVector): A vector of weights (betas) for the model.
        intercept (float): The model's intercept.

    Returns:
        float: A probability between 0 and 1.
    """
    rawPrediction = x.dot(w) + intercept #<FILL IN>

    # Bound the raw prediction value
    rawPrediction = min(rawPrediction, 20)
    rawPrediction = max(rawPrediction, -20)
    return 1/(1+exp(-rawPrediction)) #<FILL IN>

trainingPredictions = OHETrainData.map(lambda p: getP(p.features, model0.w
print trainingPredictions.take(5)

[0.30262882023911114, 0.10362661997434088, 0.2836342478387561, 0.178461
0205788012, 0.5389775379218853]
```

```
In [54]: # TEST Predicted probability (4d)
Test.assertTrue(np.allclose(trainingPredictions.sum(), 18135.4834348),
                'incorrect value for trainingPredictions')
1 test passed.
```

4e

```
In [55]: # TODO: Replace <FILL IN> with appropriate code
def evaluateResults(model, data):
    """Calculates the log loss for the data given the model.

    Args:
        model (LogisticRegressionModel): A trained logistic regression model
        data (RDD of LabeledPoint): Labels and features for each observation

    Returns:
        float: Log loss for the data.
    """
    #<FILL IN>
    return data.map(lambda p: computeLogLoss(getP(p.features, model.weights)))

logLossTrLR0 = evaluateResults(model0, OHETrainData)
print ('OHE Features Train Logloss:\n\tBaseline = {0:.3f}\n\tLogReg = {1:.3f}'.format(logLossTrBase, logLossTrLR0))

OHE Features Train Logloss:
    Baseline = 0.536
    LogReg = 0.456903
```

```
In [56]: # TEST Evaluate the model (4e)
Test.assertTrue(np.allclose(logLossTrLR0, 0.456903), 'incorrect value for logLossTrLR0')
1 test passed.
```

4f

```
In [57]: # TODO: Replace <FILL IN> with appropriate code
logLossValBase = OHEValidationData.map(lambda p: computeLogLoss(classOneFromFeatures(p.features)))

logLossValLR0 = evaluateResults(model0, OHEValidationData) #<FILL IN>
print ('OHE Features Validation Logloss:\n\tBaseline = {0:.3f}\n\tLogReg = {1:.3f}'.format(logLossValBase, logLossValLR0))

OHE Features Validation Logloss:
    Baseline = 0.528
    LogReg = 0.457
```


```
In [58]: # TEST Validation log loss (4f)
Test.assertTrue(np.allclose(logLossValBase, 0.527603), 'incorrect value for logLossValBase')
Test.assertTrue(np.allclose(logLossValLR0, 0.456957), 'incorrect value for logLossValLR0')
1 test passed.
1 test passed.
```

```
In [59]: labelsAndScores = OHEValidationData.map(lambda lp:
                                                (lp.label, getP(lp.features, m
labelsAndWeights = labelsAndScores.collect()
labelsAndWeights.sort(key=lambda (k, v): v, reverse=True)
labelsByWeight = np.array([k for (k, v) in labelsAndWeights])

length = labelsByWeight.size
truePositives = labelsByWeight.cumsum()
numPositive = truePositives[-1]
falsePositives = np.arange(1.0, length + 1, 1.) - truePositives

truePositiveRate = truePositives / numPositive
falsePositiveRate = falsePositives / (length - numPositive)

# Generate layout and plot data
fig, ax = preparePlot(np.arange(0., 1.1, 0.1), np.arange(0., 1.1, 0.1))
ax.set_xlim(-.05, 1.05), ax.set_ylim(-.05, 1.05)
ax.set_ylabel('True Positive Rate (Sensitivity)')
ax.set_xlabel('False Positive Rate (1 - Specificity)')
plt.plot(falsePositiveRate, truePositiveRate, color='#8cbfd0', linestyle='
plt.plot((0., 1.), (0., 1.), linestyle='--', color='#d6ebf2', linewidth=2.
pass
```



Part 5

5a

```
In [61]: from collections import defaultdict
import hashlib

def hashFunction(numBuckets, rawFeats, printMapping=False):
    """Calculate a feature dictionary for an observation's features based

    Note:
        Use printMapping=True for debug purposes and to better understand

    Args:
        numBuckets (int): Number of buckets to use as features.
        rawFeats (list of (int, str)): A list of features for an observati
            (featureID, value) tuples.
        printMapping (bool, optional): If true, the mappings of featureStr
            printed.

    Returns:
        dict of int to float: The keys will be integers which represent t
            features have been hashed to. The value for a given key will
            (featureID, value) tuples that have hashed to that key.
    """
    mapping = {}
    for ind, category in rawFeats:
        featureString = category + str(ind)
        mapping[featureString] = int(int(hashlib.md5(featureString).hexdigest()
    if(printMapping): print mapping
    sparseFeatures = defaultdict(float)
    for bucket in mapping.values():
        sparseFeatures[bucket] += 1.0
    return dict(sparseFeatures)
```

```
In [62]: # TODO: Replace <FILL IN> with appropriate code
# Use four buckets
sampOneFourBuckets = hashFunction(4, sampleOne, True)
sampTwoFourBuckets = hashFunction(4, sampleTwo, True)
sampThreeFourBuckets = hashFunction(4, sampleThree, True)

# Use one hundred buckets
sampOneHundredBuckets = hashFunction(100, sampleOne, True)
sampTwoHundredBuckets = hashFunction(100, sampleTwo, True)
sampThreeHundredBuckets = hashFunction(100, sampleThree, True)

print '\t\t 4 Buckets \t\t\t 100 Buckets'
print 'SampleOne:\t {0}\t\t {1}'.format(sampOneFourBuckets, sampOneHundredBuckets)
print 'SampleTwo:\t {0}\t\t {1}'.format(sampTwoFourBuckets, sampTwoHundredBuckets)
print 'SampleThree:\t {0}\t\t {1}'.format(sampThreeFourBuckets, sampThreeHundredBuckets)

{'black1': 2, 'mouse0': 3}
{'cat0': 0, 'tabby1': 0, 'mouse2': 2}
{'bear0': 0, 'black1': 2, 'salmon2': 1}
{'black1': 14, 'mouse0': 31}
{'cat0': 40, 'tabby1': 16, 'mouse2': 62}
{'bear0': 72, 'black1': 14, 'salmon2': 5}

4 Buckets                                100 Buckets
SampleOne:      {2: 1.0, 3: 1.0}          {14: 1.0, 31: 1.0}
SampleTwo:      {0: 2.0, 2: 1.0}          {40: 1.0, 16: 1.0, 62:
1.0}
SampleThree:    {0: 1.0, 1: 1.0, 2: 1.0}   {72: 1.0, 5: 1.0, 14:
1.0}
```

```
In [63]: # TEST Hash function (5a)
Test.assertEquals(sampOneFourBuckets, {2: 1.0, 3: 1.0}, 'incorrect value f
Test.assertEquals(sampThreeHundredBuckets, {72: 1.0, 5: 1.0, 14: 1.0},
'incorrect value for sampThreeHundredBuckets')

1 test passed.
1 test passed.
```

5b

[illegible]


```

In [65]: # TEST Creating hashed features (5b)
hashTrainDataFeatureSum = sum(hashTrainData
                                .map(lambda lp: len(lp.features.indices))
                                .take(20))
hashTrainDataLabelSum = sum(hashTrainData
                              .map(lambda lp: lp.label)
                              .take(100))
hashValidationDataFeatureSum = sum(hashValidationData
                                     .map(lambda lp: len(lp.features.indices))
                                     .take(20))
hashValidationDataLabelSum = sum(hashValidationData
                                   .map(lambda lp: lp.label)
                                   .take(100))
hashTestDataFeatureSum = sum(hashTestData
                              .map(lambda lp: len(lp.features.indices))
                              .take(20))
hashTestDataLabelSum = sum(hashTestData
                            .map(lambda lp: lp.label)
                            .take(100))

Test.assertEquals(hashTrainDataFeatureSum, 772, 'incorrect number of featu
Test.assertEquals(hashTrainDataLabelSum, 24.0, 'incorrect labels in hashTr
Test.assertEquals(hashValidationDataFeatureSum, 776,
                  'incorrect number of features in hashValidationData')
Test.assertEquals(hashValidationDataLabelSum, 16.0, 'incorrect labels in h
Test.assertEquals(hashTestDataFeatureSum, 774, 'incorrect number of featur
Test.assertEquals(hashTestDataLabelSum, 22.0, 'incorrect labels in hashTe

1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.

```

5c

```
In [66]: # TODO: Replace <FILL IN> with appropriate code
def computeSparsity(data, d, n):
    """Calculates the average sparsity for the features in an RDD of LabeledPoint.

    Args:
        data (RDD of LabeledPoint): The LabeledPoints to use in the sparsity calculation.
        d (int): The total number of features.
        n (int): The number of observations in the RDD.

    Returns:
        float: The average of the ratio of features in a point to total features.
    """
    #<FILL IN>
    return data.map(lambda p: 1.0*len(p.features.indices)/d).mean()

averageSparsityHash = computeSparsity(hashTrainData, numBucketsCTR, nTrain)
averageSparsityOHE = computeSparsity(OHETrainData, numCtrOHEFeats, nTrain)

print 'Average OHE Sparsity: {0:.7e}'.format(averageSparsityOHE)
print 'Average Hash Sparsity: {0:.7e}'.format(averageSparsityHash)

Average OHE Sparsity: 1.6717677e-04
Average Hash Sparsity: 1.1805561e-03
```

```
In [67]: # TEST Sparsity (5c)
Test.assertTrue(np.allclose(averageSparsityOHE, 1.6717677e-04),
                'incorrect value for averageSparsityOHE')
Test.assertTrue(np.allclose(averageSparsityHash, 1.1805561e-03),
                'incorrect value for averageSparsityHash')

1 test passed.
1 test passed.
```

5d

```
In [68]: numIters = 500
regType = 'l2'
includeIntercept = True

# Initialize variables using values from initial model training
bestModel = None
bestTestLoss = 1e10
```

```
In [69]: # TODO: Replace <FILL IN> with appropriate code
stepSizes = [1, 10] #<FILL IN>
regParams = [1e-6, 1e-3] #<FILL IN>
for stepSize in stepSizes:
    for regParam in regParams:
        model = (LogisticRegressionWithSGD
                  .train(hashTrainData, numIters, stepSize, regParam=regParam,
                        intercept=includeIntercept))
        logLossVa = evaluateResults(model, hashValidationData)
        print ('\tstepSize = {0:.1f}, regParam = {1:.0e}: logloss = {2:.6f}'.format(stepSize, regParam, logLossVa))
        if (logLossVa < bestLogLoss):
            bestModel = model
            bestLogLoss = logLossVa

print ('Hashed Features Validation Logloss:\n\tBaseline = {0:.6f}\n\tLogReg = {1:.6f}'.format(logLossVa, bestLogLoss))

stepSize = 1.0, regParam = 1e-06: logloss = 0.474694
stepSize = 1.0, regParam = 1e-03: logloss = 0.474999
stepSize = 10.0, regParam = 1e-06: logloss = 0.449679
stepSize = 10.0, regParam = 1e-03: logloss = 0.451841
Hashed Features Validation Logloss:
Baseline = 0.527603
LogReg = 0.449679
```

```
In [70]: # TEST Logistic model with hashed features (5d)
Test result True (expected logloss (bestTestLogLoss = 0.4491692609) - incorrect value of
1 test failed. incorrect value for bestLogLoss
```

Visualization

```
In [71]: from matplotlib.colors import LinearSegmentedColormap


# Saved parameters and results. Eliminate the time required to run 36 mod
stepSizes = [3, 6, 9, 12, 15, 18]
regParams = [1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2]
logLoss = np.array([[ 0.45808431,  0.45808493,  0.45809113,  0.45815333,
                      [ 0.45188196,  0.45188306,  0.4518941,  0.4520051,
                      [ 0.44886478,  0.44886613,  0.44887974,  0.44902096,
                      [ 0.44706645,  0.4470698,  0.44708102,  0.44724251,
                      [ 0.44588848,  0.44589365,  0.44590568,  0.44606631,
                      [ 0.44508948,  0.44509474,  0.44510274,  0.44525007,

numRows, numCols = len(stepSizes), len(regParams)
logLoss = np.array(logLoss)
logLoss.shape = (numRows, numCols)

fig, ax = preparePlot(np.arange(0, numCols, 1), np.arange(0, numRows, 1),
                      hideLabels=True, gridWidth=0.)
ax.set_xticklabels(regParams), ax.set_yticklabels(stepSizes)
ax.set_xlabel('Regularization Parameter'), ax.set_ylabel('Step Size')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055'],
image = plt.imshow(logLoss, interpolation='nearest', aspect='auto',
                    cmap = colors)

#pass
plt.show()
```



5e

```
In [72]: # TODO: Replace <FILL IN> with appropriate code
# Log loss for the best model from (5d)
best = (LogisticRegressionWithSGD.train(hashTrainData,
                                         iterations = numIters,
                                         step = 17,
                                         regParam=1e-05,
                                         regType=regType,
                                         intercept=includeIntercept))

logLossTest = evaluateResults(best, hashTestData) #<FILL IN>

# Log loss for the baseline model
logLossTestBaseline = hashTestData.map(lambda p: computeLogLoss(classOneFr

print ('Hashed Features Test Log Loss:\n\tBaseline = {0:.6f}\n\tLogReg = {
      .format(logLossTestBaseline, logLossTest))
```

```
Hashed Features Test Log Loss:
    Baseline = 0.537438
    LogReg = 0.453568
```

```
In [73]: # TEST Evaluate on the test set (5e)
Test.assertTrue(np.allclose(logLossTestBaseline, 0.537438),
                'incorrect value for logLossTestBaseline')
Test.assertTrue(np.allclose(logLossTest, 0.453568), 'incorrect value for
1 test passed.
1 test failed. incorrect value for logLossTest
```

```
In [ ]:
```